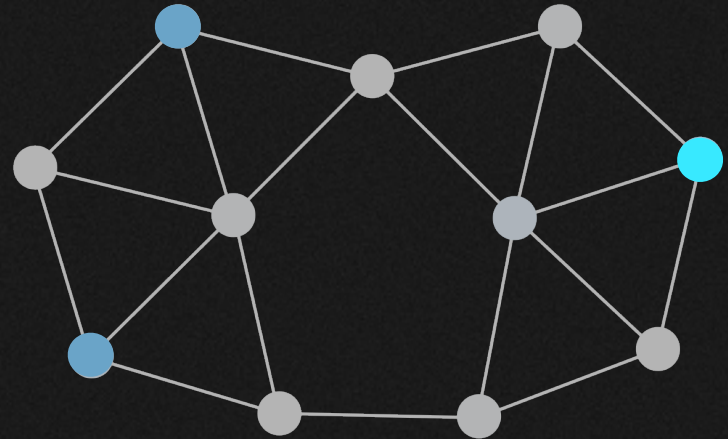


Distributed Systems

HELLO WORLD



Introduction to MPI



Message Passing Interface (MPI)

- Standard for Message Passing Computation: a collection of processes communicating with messages.
- Cross Platform (desktop to supercomputer)
- Language Agnostic (C, C++, FORTRAN...)

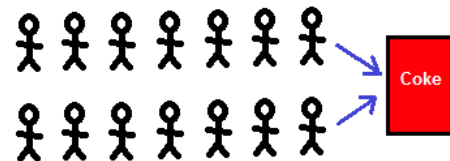


Message Passing Interface (MPI)

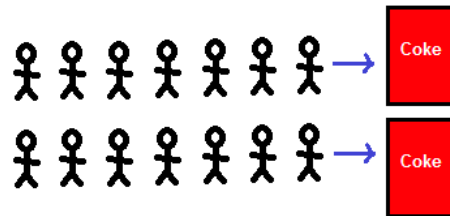
- MPI is a library, not a language.
- Specifies: the names, calling sequences and results of functions or subroutines to be called from C/C++ or Fortran programs, and the classes and methods that make up the MPI C++ library.
- Programs that users write are compiled with ordinary compilers and linked with the MPI library.

Parallelism

- Task parallelism: the work of a global problem can be divided into a number of independent tasks, which rarely need to synchronize.
- E.g. Monte Carlo simulations or numerical integration



Concurrent: 2 queues, 1 vending machine



Parallel: 2 queues, 2 vending machines



Disadvantages/Issues

- No free lunch - can't just “turn on” parallel
- Parallel programming requires work
 - Code modification - always
 - Algorithm modification - often
 - New sneaky bugs - you bet
- Speedup limited by many factors

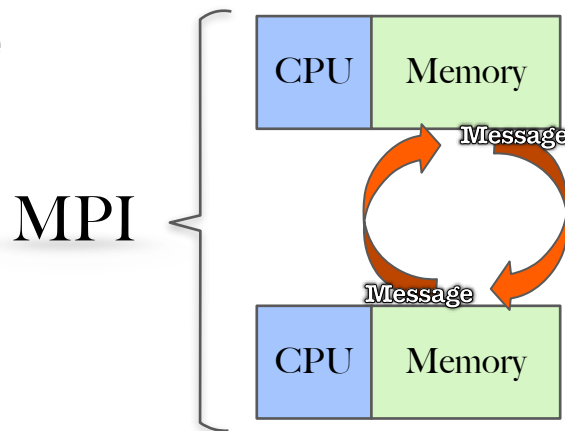


Computer Architecture

- As you consider parallel programming understanding the underlying architecture is important
- Performance is affected by hardware configuration
 - Memory or CPU architecture
 - Numbers of cores/processor
 - Network speed and architecture

MPI is a library

- MPI is a library specification for the message passing interface, proposed as a standard.
 - independent of hardware;
 - not a language or compiler specification;
 - not a specific implementation or product.
- A message passing standard for portability and ease-of-use.
- Designed for high performance.



MPI Program - Basics

Include MPI Header File

Start of Program
(Non-interacting Code)

Initialize MPI

Run Parallel Code &
Pass Messages

End MPI Environment

(Non-interacting Code)

End of Program

```
#include <mpi.h>
```

```
int main(int argc, char *argv[])  
{
```

```
    MPI_Init(&argc, &argv);
```

```
    .
```

```
    .    // Run parallel code
```

```
    .
```

```
    MPI_Finalize();    //End MPI Environment
```

```
    return 0;
```

```
}
```



MPI Common Terms

Communicator

- A group of MPI processes with a name (context).
- Processes identified by rank (only meaningful within particular communicator).
- Mechanism to identify subset of processes.

`MPI_COMM_WORLD`

- Predefined default communicator
- Contains all of the MPI Processes
- Size = N_{procs}



MPI: Most important functions

<code>MPI_Init</code>	Initiate an MPI computation
<code>MPI_Finalize</code>	terminate the MPI computation and clean up
<code>MPI_Comm_size</code>	how many processes participate in a given MPI communicator?
<code>MPI_Comm_rank</code>	which one am I? (Number between 0 and size-1)
<code>MPI_Send</code>	send a message to a particular process within an MPI communicator
<code>MPI_Recv</code>	receive a message from a particular process within an MPI communicator
<code>MPI_Bcast</code>	send a message to all processes within an MPI communicator
<code>MPI_Reduce</code>	collect overall data and reduce to a single datum using a specified operation
<code>MPI_Gather</code>	collect overall data from all processes within an MPI communicator
<code>MPI_Scatter</code>	divide a large amount of data among all processes within an MPI communicator



Output

- Output to screen is not ordered
- All processes are trying to write to screen simultaneously.
- The OS opts for an ordering.
- For organized output, starting from the first process, we can rewrite our program (E.g. Using MPI_Barrier)

Native install

Linux



[Read Document on Moodle](#) → MPI_setup - Linux.pdf

```
sudo apt-get install build-essential libopenmpi-dev openmpi-bin openmpi-doc
```

(or whatever package manager you use)

Mac



Read Document on Moodle → MPI_setup - Mac.pdf

<http://brew.sh>

In a terminal: `brew install open-mpi`

Caveats:

1. *"There are not enough slots available in the system to satisfy the x slots"*

Add the following argument to your `mpirun` call: `--oversubscribe`

Ex: `mpirun -n 4 --oversubscribe ./myfile.exe`

2. You might need to elevate (sudo) the *mpirun* command

Windows



[Read Document on Moodle](#)



MPI_Setup - Windows.pdf

- ☒ Install VisualStudio 2022 with C++ workspace
- ☒ Install ms-mpi
- ☒ Configure VS to use MSMPI library

Virtualbox

Only if already familiar with Linux !

Create a VM

Download & install VirtualBox.

Install Ubuntu Server LTS in a VM.

Allow SSH access to VM.

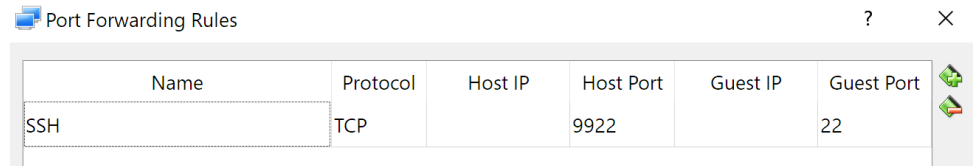
Use putty to connect to VM.

127.0.0.1:9922

Install software:

```
sudo apt-get install build-essential libopenmpi-dev openmpi-bin openmpi-doc
```

Network->Advanced->Port Forwarding



Name	Protocol	Host IP	Host Port	Guest IP	Guest Port
SSH	TCP		9922		22



Tweak Linux

Add the following lines to the following two files:

```
sudo nano /etc/openmpi/openmpi-mca-params.conf  
    btl_base_warn_component_unused = 0
```

```
sudo nano /etc/openmpi/openmpi-default-hostfile  
    localhost slots=4
```



First App



Important

Complete the instructions for installation and setup for your own OS on Moodle before attempting to run the following programs.



Hello World

Create a hello world file: "hello.cpp".

Add some code. →

```
1 //This is a comment
2 #include <iostream> //preprocessor - looks for file, and puts it here
3
4 //The main() function is a special function, it's the beginning
5 int main(int argc, char **argv) {
6     //print a message to the screen
7     std::cout << "Hello world!" << std::endl;
8     return 0; //return
9 }
```

Compile: mpic++ ./hello.cpp -o ./hello
-or-
Build using IDE

Run: mpirun ./hello
-or-
mpiexec ./hello

Note: *mpirun*, *mpiexec*, and *orterun* are all synonyms for each other. Using any of the names will produce the same behavior.



What is mpirun?

mpirun [options] **<program>** [<args>]

- A command which controls several aspects of program execution in Open MPI.
- Uses the Open Run-Time Environment (ORTE) to launch jobs.
- If you are running under distributed resource manager software, such as Sun Grid Engine or PBS, ORTE launches the resource manager for you.

More info:

<http://www.sciencedirect.com/science/article/pii/S0167739X07000507>

https://docs.oracle.com/cd/E19356-01/820-3176-10/ExecutingPrograms.html#50413574_62903

Note: *mpirun*, *mpiexec*, and *orterun* are all synonyms for each other. Using any of the names will produce the same behavior.



What is mpiexec?

mpiexec [options] **<program>** [<args>]

- A command which controls several aspects of program execution in Open MPI.
- Uses the Open Run-Time Environment (ORTE) to launch jobs.
- If you are running under distributed resource manager software, such as Sun Grid Engine or PBS, ORTE launches the resource manager for you.

More info:

<http://www.sciencedirect.com/science/article/pii/S0167739X07000507>

https://docs.oracle.com/cd/E19356-01/820-3176-10/ExecutingPrograms.html#50413574_62903

Note: *mpirun*, *mpiexec*, and *orterun* are all synonyms for each other. Using any of the names will produce the same behavior.



What is **orterun**?

orterun [options] **<program>** [<args>]

- A command which controls several aspects of program execution in Open MPI.
- Uses the Open Run-Time Environment (ORTE) to launch jobs.
- If you are running under distributed resource manager software, such as Sun Grid Engine or PBS, ORTE launches the resource manager for you.

More info:

<http://www.sciencedirect.com/science/article/pii/S0167739X07000507>

https://docs.oracle.com/cd/E19356-01/820-3176-10/ExecutingPrograms.html#50413574_62903



Compilation and Execution Depends on your OS!

To compile your file: `mpic++ NameOfFile.cpp -o NameOfFile`

This will create an executable file with that name, that you can then run using the mpirun command:

```
mpirun -np (number of nodes) --oversubscribe (name of file)
```



Compilation and Execution Windows + VS

To compile your file: `Build Solution (ctrl+shift+B)`

This will create an executable file with the same name as the original .cpp file, that you can then run using the mpiexec command:

```
mpiexec -np (number of nodes) --oversubscribe (name of file)
```



Compilation and Execution OS + Xcode

To compile your file: Build Solution (⌘+B)

This will create an executable file with the same name as the original .cpp file, that you can then run using the mpirun command:

```
mpirun -np (number of nodes) --oversubscribe (name of file)
```



Open MPI App

Now that you've run a hello world using c++, try importing the MPI library and running using more than one node.


Create a new file: "hello2.cpp".

Compile: `mpic++ ./hello2.cpp -o ./hello2`

Run: `mpirun -np 4 --oversubscribe ./hello2`

```
1 #include <mpi.h>
2 #include <stdio.h>
3
4 int main(int argc, char** argv) {
5     MPI_Init(NULL, NULL);
6     int rank;
7     int world;
8     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
9     MPI_Comm_size(MPI_COMM_WORLD, &world);
10    printf("Hello: rank %d, world: %d\n", rank, world);
11    MPI_Finalize();
12 }
```

Trouble?



Send() fails?

tcp_peer_send_blocking: send() to socket 12 failed

```
netstat -rn
```

```
route print
```

Should be the same IP as the first one in:

```
mpirun -report-uri - -n 4 ./a.exe
```

So disable NICs to make sure MPI uses your correct IP.



OpenCL errors when running in Xcode

- You may ignore these, they only appear when running directly via the IDE.

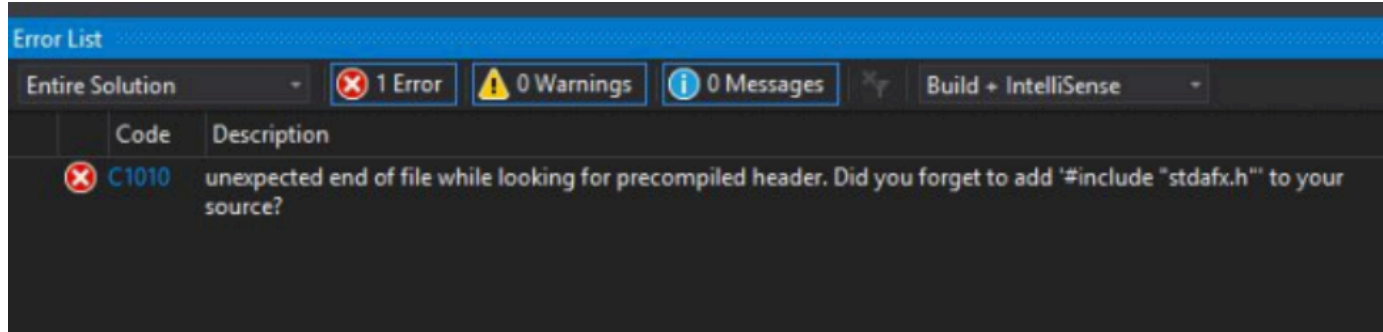
```
2022-09-05 16:56:12.363081+0100 orterun[636:6500280] [CL_INVALID_OPERATION] : OpenCL Error :  
Failed to retrieve device information! Invalid enumerated value!  
  
2022-09-05 16:56:12.364089+0100 orterun[636:6500280] [CL_INVALID_OPERATION] : OpenCL Error :  
Failed to retrieve device information! Invalid enumerated value!  
  
2022-09-05 16:56:12.364143+0100 orterun[636:6500280] [CL_INVALID_OPERATION] : OpenCL Error :  
Failed to retrieve device information! Invalid enumerated value!
```




Linking Error Messages (Windows)

- If you get any linking error messages, it is most likely you're building for 64 bits yet specifying 32 bits linking libraries.
- Ensure you are using the correct linking libraries for your architecture
- You can check your architecture by going to Windows > Settings > System > About > System type, this should tell you your processor type.

Unexpected end of file C1010



- Go to Property Pages – C/C++ Precompiled Header then click Precompiled Header and select Not Using Precompiled Headers.
- Note: Precompiled headers can be turned off when you create a new project!



Add exceptions for...

Firewall

AV

LabWork - Week 1



Lab Work - Week 1

- Complete installation and run your first two Hello World programs.
 - ➡ One using only C++
 - ➡ One using C++ and MPI
- Start MPI_Notes tutorials